

Using Deep Reinforcement Learning for Autonomous Cars in 2D and 3D Simulated Environments

ECE 510 - Deep Learning Theory and Practice
Final Project - Spring 2019

Husnu Melih Erdogan

Mohammed Abidalrekab

Abdirahman Alasow



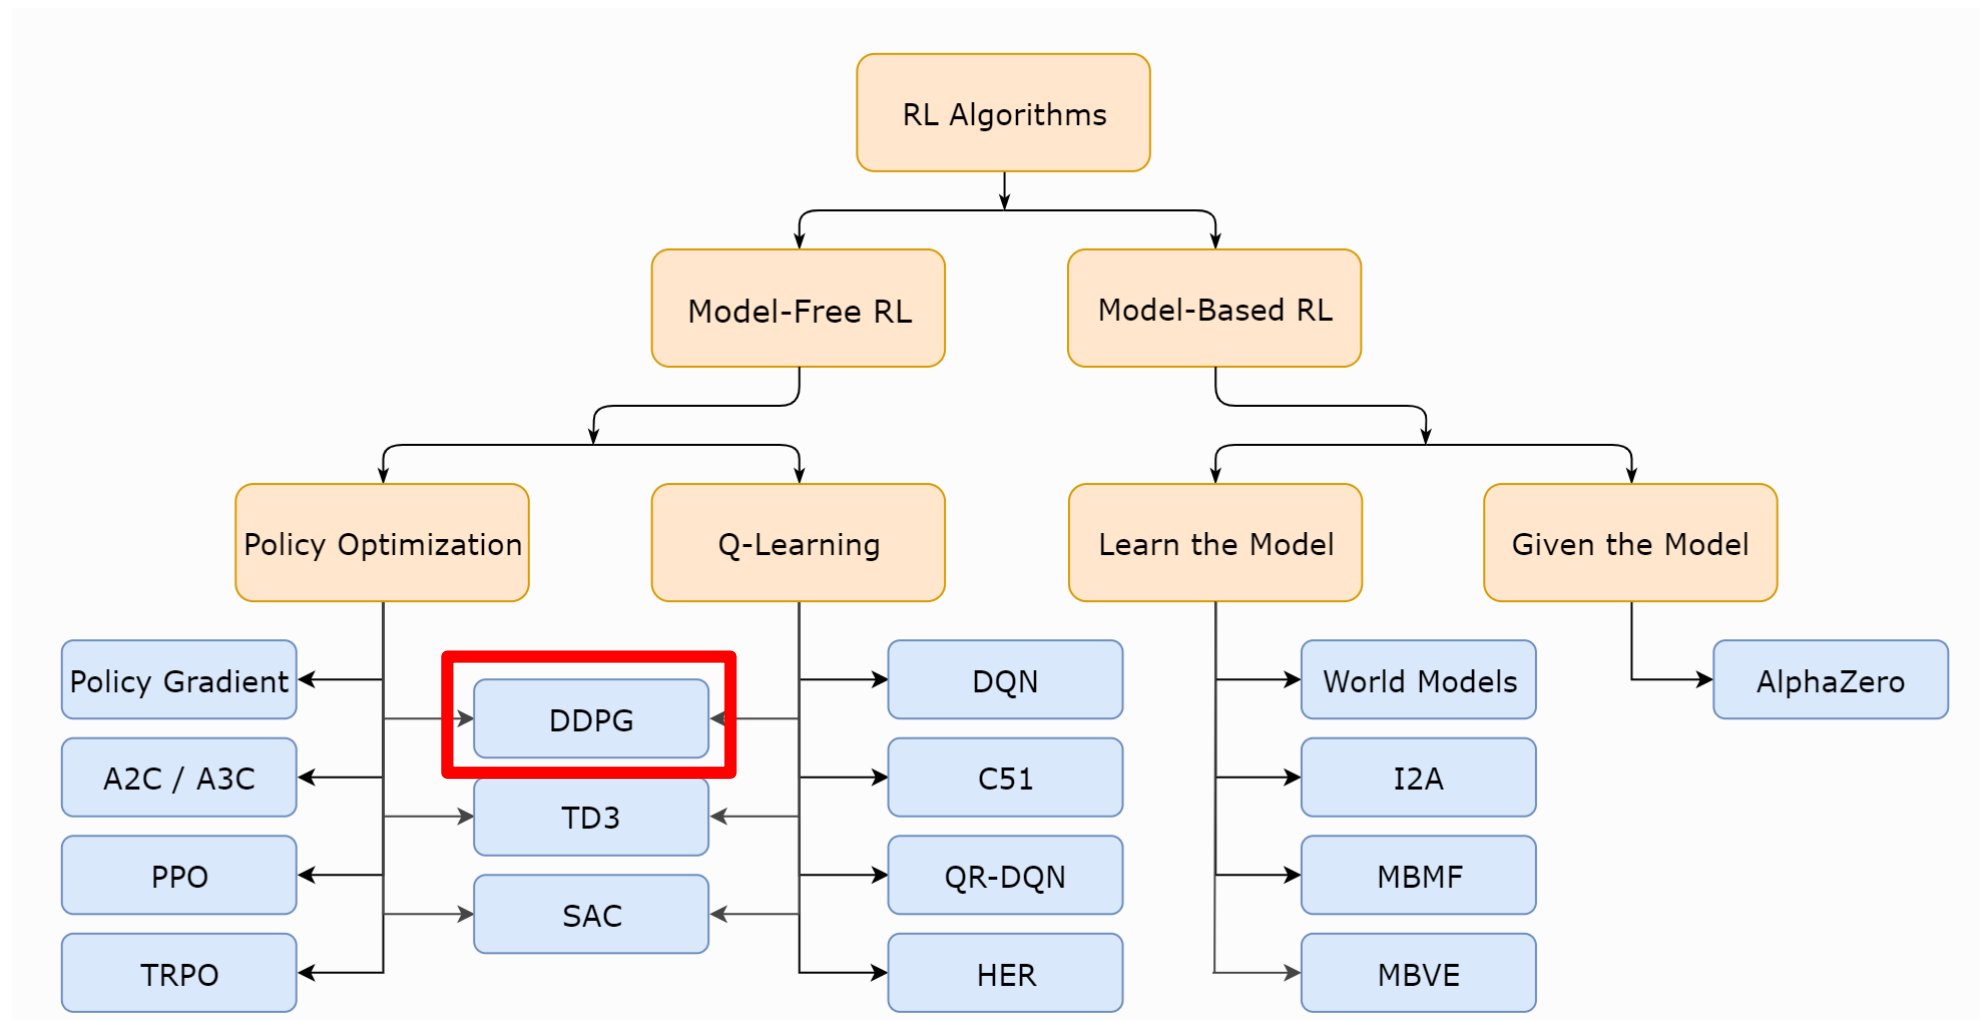
Goal

- Our main goal in this project is discovering and learning more about Deep Reinforcement Learning for autonomous vehicle navigation as we build and test a self driving car application in a simulated environment.
- We compared several different Deep Reinforcement methods and used them with several different simulations.

Papers

- Continuous Control With Deep Reinforcement Learning
 - Timothy P. Lillicrap* , Jonathan J. Hunt* , Alexander Pritzel, Nicolas Heess, Tom Erez, Yuval Tassa, David Silver & Daan Wierstra
 - Google Deepmind – 2016
- Continuous Control Automated Lane Change Behavior Based on Deep Deterministic Policy Gradient Algorithm
 - Pin Wang, Hanhan Li, Ching-Yao Chan
 - IEEE IVS, 2019
- Soft Actor-Critic – Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor
 - Thomas Haarnoja, Aurick Zhou, Pieter Abbel, Sergey Levine
 - ICML, 2018

Taxonomy of Reinforcement Learning Algorithms



Simulations


- TORCS

- 3D Car Simulator
- 29 States (speed, wheelspin, rpm, track etc.)
- 3 Actions (Steering, Acceleration, Brake)




- Flatlands

- 2D Car Simulator
- 3 States (Distances from the track)
- 2 Actions (Steering, Acceleration, Brake)




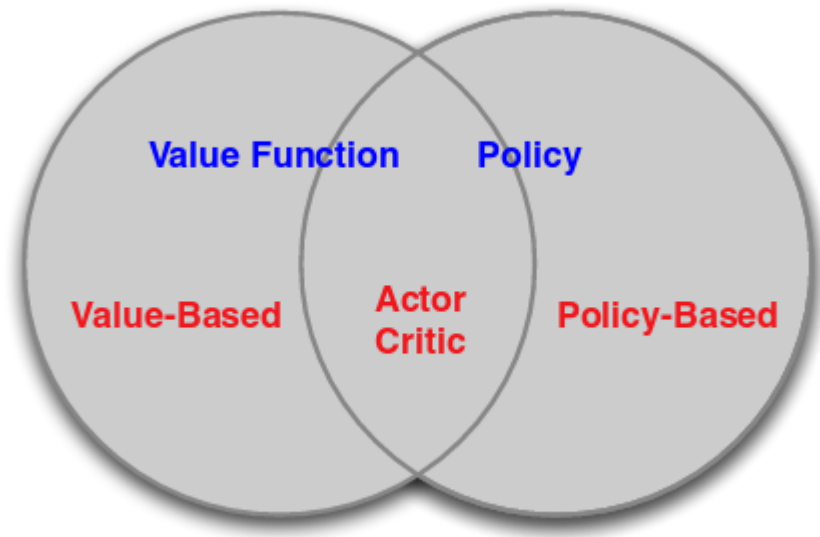
- Pendulum

- 2D Pendulum Simulator
- 3 States (position, velocity, angle)
- 1 Actions (Joint effort)



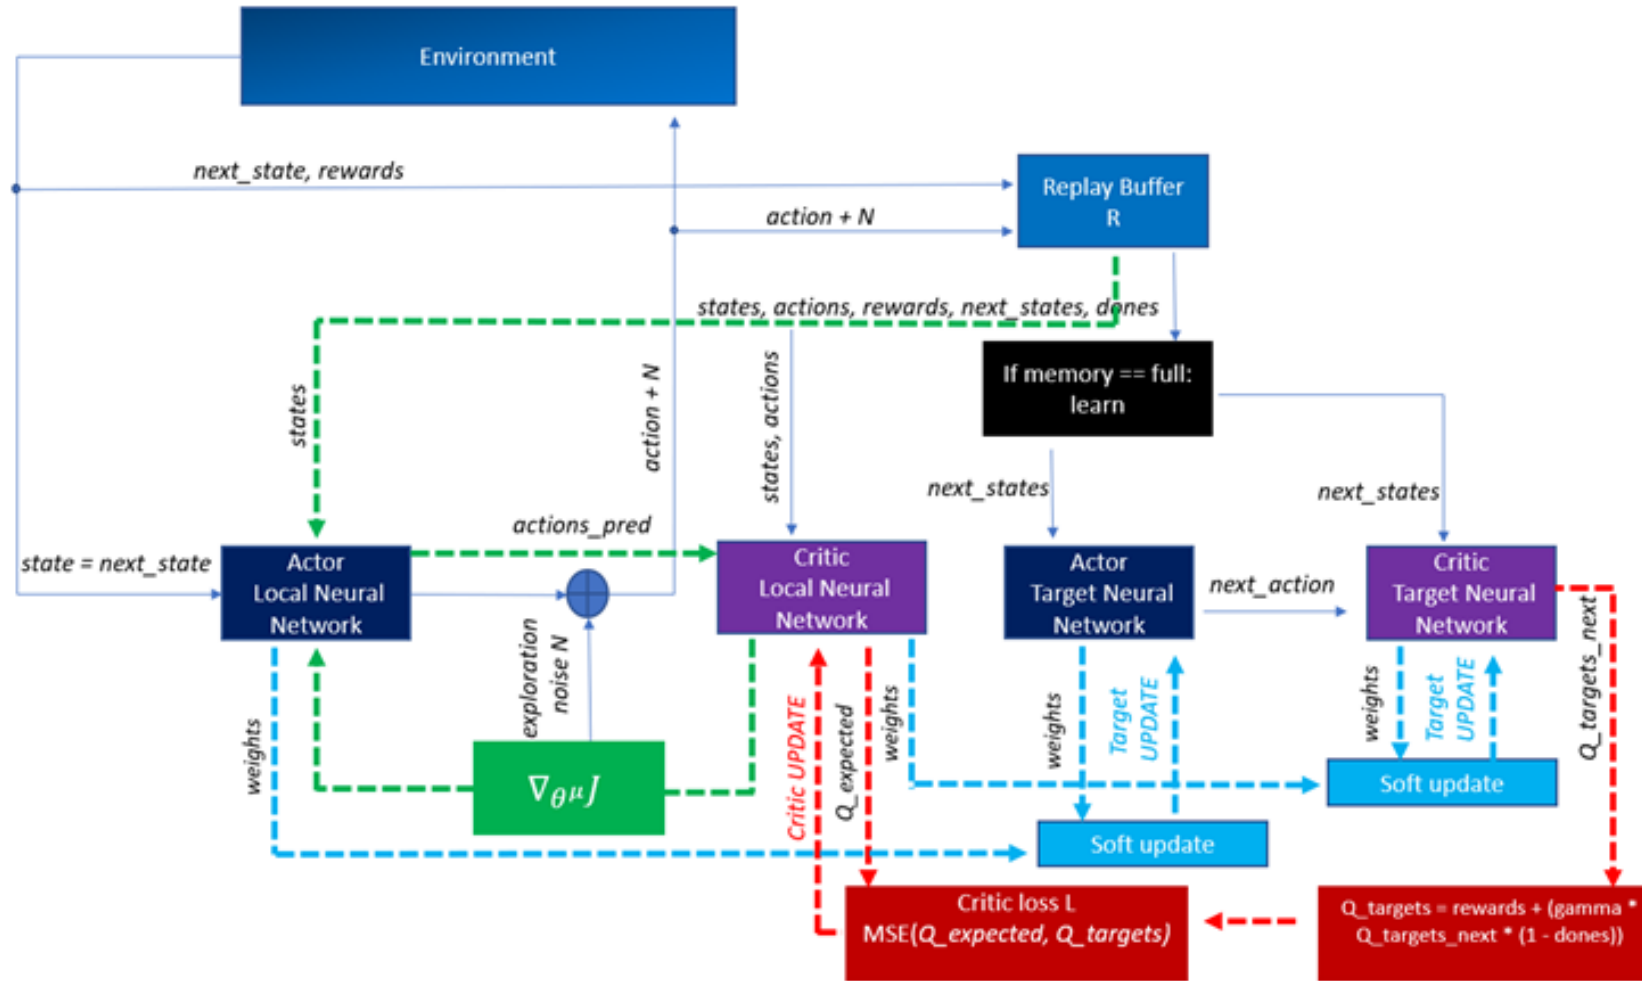
DDPG

- DDPG - Deep Deterministic Policy Gradient

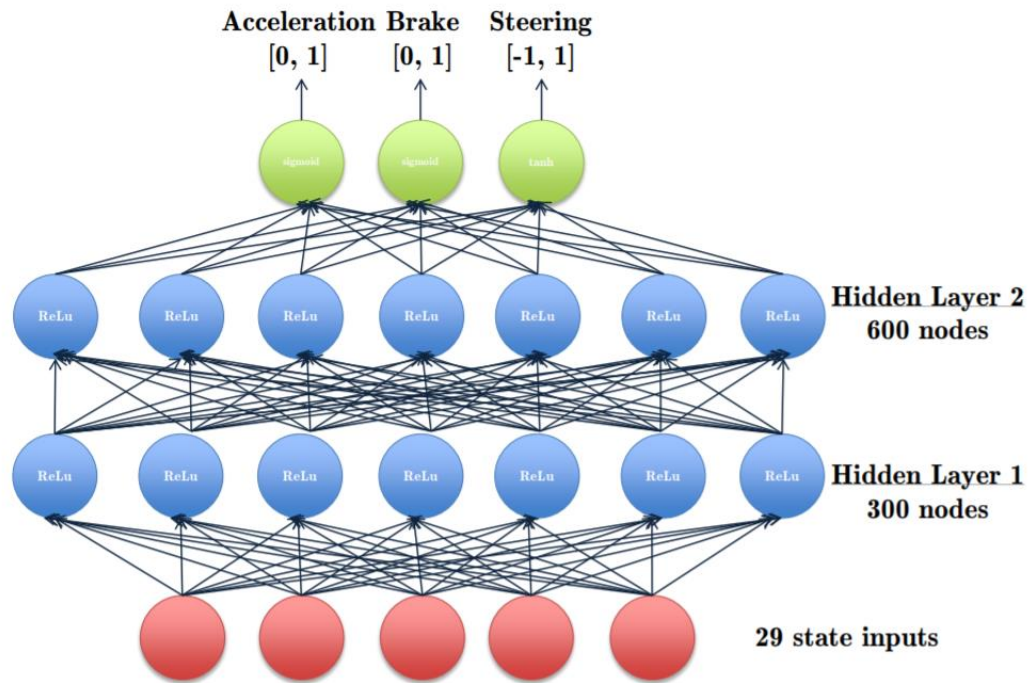


$$\nabla_{\theta} J(\theta) \sim \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) A(s_t, a_t)$$


DDPG



DDPG



(a) Actor network with two ReLu activated hidden layers



(b) Critic network with ReLu and linear activated hidden layers

DDPG

Algorithm 1 DDPG algorithm

Randomly initialize critic network $Q(s, a|\theta^Q)$ and actor $\mu(s|\theta^\mu)$ with weights θ^Q and θ^μ .

Initialize target network Q' and μ' with weights $\theta^{Q'} \leftarrow \theta^Q$, $\theta^{\mu'} \leftarrow \theta^\mu$

Initialize replay buffer R

for episode = 1, M **do**

 Initialize a random process \mathcal{N} for action exploration

 Receive initial observation state s_1

for t = 1, T **do**

 Select action $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$ according to the current policy and exploration noise

 Execute action a_t and observe reward r_t and observe new state s_{t+1}

 Store transition (s_t, a_t, r_t, s_{t+1}) in R

 Sample a random minibatch of N transitions (s_i, a_i, r_i, s_{i+1}) from R

 Set $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'})|\theta^{Q'})$

 Update critic by minimizing the loss: $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$

 Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

 Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$


$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$

end for

end for

Experiments - Pendulum

- Beginning




- We did this experiment to understand and prove the algorithm

- Training time:


- 50 episodes (~5 mins)

- Result:

- Shows that it learns the desired behavior
- Able to keep the pendulum up




Experiment - Pendulum




Experiments - TORCS

- Beginning




- Result





- We did this experiment to test the DDPG method for our application idea
- Training time:
 - 100 episodes (~30 mins)
 - Training this car simulation using CPU only requires ~2-3 days
- Results:
 - Shows that it is on the way of learning the desired behavior
 - Able to make turns and keep the car straight at a certain speed.

Experiment - TORCS



>10 hours


Experiment - Flatlands




- We did this experiment to test our work
- Training time:
 - 10 episodes (~15 mins)
 - 20 episodes (~30 mins)
- Tests:
 - Tested over 20 different reward functions
 - Tested different noise levels
 - Different buffer sizes

Final Result

$\text{Reward_funtions} = 30 * \text{abs}(c_speed) - 10 * \text{abs}(r_speed) - 5 * (\text{abs}(\text{math.sqrt}(\text{distance_x} ** 2 + \text{distance_y} ** 2))) - 2 * \text{head_x}$



Final Result



Future Work

- Testing different reward functions
- Testing the effects of different step size
- Using different RL algorithms to solve the same problems
 - PPO
 - A2C
 - A3C
- Testing our work on different benchmarks

Conclusion

- We were able to train our agent (self-driving car) to learn how to stay closer to the track and move
- We had an opportunity to use the tools and knowledge that we learned in this course such as Pytorch, Deep Reinforcement Learning, MLP.
- Finding a good simulator that satisfies our needs for this project was a challenge.
- We have tested many different simulations built for OpenAI gym and Unity



Thanks for listening!